

ICEBERG

Phase Execution Playbook

Seventeen phases. Five audit checkpoints. One repeatable ritual.

This playbook consolidates PLAN.md, the AI Session Prompt Library (98-governance/11), and the Founder Execution Guide into one printable execution companion: a guide, a detailed prompt, a model assignment, and exit checks for every phase - plus a guided connectivity audit after each phase cluster.

Precedence: this PDF is a convenience view (Tier 5). PLAN.md, ARCHITECTURE.md, and the governance documents are the binding law. If this PDF ever disagrees with them, they win and the discrepancy is an ERRATA report.

Repository	D:\Amazon Product Reasearch\iceberg (code) - iceberg-docs (law)
Status at issue	Phase 0 in progress - T0.1 complete (module boundaries verified by violation test)
Issued	2026-07-13 - v1.0

1 - How to Use This Playbook

Every implementation session follows the same ritual. Print this page.

- 1 Open PLAN.md -> current phase -> pick ONE task. Check PLAN-VARIANCE.md and 98-governance/ERRATA.md for anything touching your files.
- 2 Start a fresh AI session. Paste the Standing Header (below), then this playbook's phase prompt, then the task ID.
- 3 Demand the READ-BACK: one paragraph stating what will be done and which documents bind it. Verify it against the spec YOURSELF. Wrong read-back = wrong day.
- 4 Implement -> test (including the NEGATIVE fixtures - the violation that must fail) -> commit small and cited.
- 5 Review in a FRESH session (prompt P4) - the writer never reviews its own work. Demo the acceptance criterion to yourself, literally.
- 6 Update task status, variance log, session note. Write tomorrow's ONE task on a sticky note.

The Standing Header (paste verbatim at the top of EVERY session)

STANDING HEADER - from 98-governance/11-AI-SESSION-PROMPT-LIBRARY.md

You are working in the Iceberg repository. Before doing anything:

1. Read PLAN.md Section 1 (rules) and Section 7 (playbook), then 98-governance/04-AI-CODING-PLAYBOOK.md.
2. Read the phase/task sections named below, then every ARCHITECTURE.md and design/intelligence section they cite. Check PLAN-VARIANCE.md and 98-governance/ERRATA.md for anything touching your files.
3. Produce the read-back: one paragraph stating what you will do, which task IDs it covers, and which documents bind it. No code before it.

Hard rules: the Ten Prohibitions (PLAN 1.4) are absolute. Never modify T0-T3 documents, Canon config, engine functions, or prompt templates outside an explicitly scoped lane. If any spec is ambiguous, STOP and file a design-defect question (quote both readings) instead of guessing. Stay in one lane per session (engine | agent-template | UX | infra | docs). End with: tests green, DoD evidence links, variance/runbook updates, 3-6 line session note.

2 - Master Timeline

PLAN.md's calendar assumed five engineers plus AI agents. As a solo founder with AI agents, apply the sanctioned stretch of 1.25-1.75x on engineering-heavy phases. **The order is law; the calendar is yours.** Log stretches in PLAN-VARIANCE.md without guilt - cutting order or gates is the only real failure.

Phase	What	Plan wks	Solo est.	Marker
P0	Repository & Foundation	1-2	Wk 1-2	in progress
P1	Identity, Workspace & Tenancy	2-4	Wk 2-5	
P2	Evidence Store & Provenance	4-6	Wk 5-7	
P3	Canon Engines	6-9	Wk 7-11	CHECKPOINT ALPHA
P4	Import Pipeline	9-13	Wk 11-17	Milestone A (w/ P5)
P5	Onboarding & App Shell (parallel w/ P4)	9-12	Wk 11-16	
P6	Research Pipeline & Object Model	13-17	Wk 17-22	
P7	Discovery Sessions	17-20	Wk 22-26	BRAVO = Milestone B
P8	Artifact Store & Claim Checker (par. w/ P7)	17-21	Wk 24-28	
P9	AI Runtime & Model Gateway	21-25	Wk 28-33	CHECKPOINT CHARLIE
P10	The Four Analyst Agents	25-33	Wk 33-44	
P11	Committee Protocol & Verdict	33-38	Wk 44-51	DELTA = Milestone C
P12	Committee Conversation (par. w/ P13)	38-41	Wk 51-55	
P13	Reports, Handoff & Sharing	38-42	Wk 51-56	
P14	Hardening & Private Beta	42-46	Wk 56-61	Milestone D
P15	Commercialization	46-49	Wk 61-64	
P16	GA Launch & Operate	49-52	Wk 64-67	ECHO -> Milestone E

Parallel track (never on this table because it never waits): PM-1 concierge verdicts run alongside P0-P3. Ten paid manual verdicts before you believe anything else. First outreach is part of Phase 0's Definition of Done.

3 - Model Selection Guide

Assignments follow one principle: **correctness-critical lanes get the deepest reasoning model; routine lanes get the fast one; review always happens in a fresh session, and for the four correctness-critical modules, preferably in a different model family** (the Devil's-Advocate principle applied to your own workflow - different families have different blind spots).

Role	Model	Use it for
Deep implementer	Claude Opus 4.8	Engines (P3), Claim Checker (P8), AI runtime (P9), committee protocol (P11), RLS/security (P1), billing money paths (P15), anything touching a Prohibition
Frontier reasoning	Claude Fable 5	Checkpoint connectivity audits, agent prompt-template design (P10-P11), the hardest debugging sessions (third attempt after two bounces)
Default implementer	Claude Sonnet 5	UI phases (P5, P6, P7, P13), parsers and fixtures (P4), tests, migrations, infra plumbing, docs - the daily workhorse
Quick chores	Claude Haiku 4.5	Commit messages, session notes, small refactors, log/format cleanups
Second family	GPT-class (latest)	Independent review of engines, identity, billing, claim-checker PRs; attack-vector brainstorming for the security checklist
Researcher	Gemini-class (latest)	Vendor docs digestion, provider-terms checks, Amazon policy updates - never implementation

Three unbreakable session rules

- The author never reviews its own work.** Review = fresh session, prompt P4; for engines/identity/billing/claim-checker also run a second-family review.
- One lane per session** (engine | agent-template | UX | infra | docs). Mixed lanes are rejected at review.
- Two bounces = change approach, not effort.** Third attempt goes to a fresh session on a stronger model with a written list of three DIFFERENT hypotheses.

Note: these assignments are workflow guidance (Tier 5), not architecture. The product's own runtime model routing (capability classes, DA on a second family) is fixed by 07-intelligence and arrives in Phase 9 - do not confuse the two.

4 - Phase Guides & Prompts

One block per phase, in execution order, with the five audit checkpoints interleaved exactly where PLAN.md places them. Prompts assume the Standing Header is already pasted. Replace <T0.x>-style slots with the specific task ID.

P0 - Repository & Foundation

Window Wk 1-2
Model Sonnet 5 (Opus 4.8 for the import-law lint)
Status IN PROGRESS - T0.1 done: boundaries proven by violation test

Objective. Enforcement before code: monorepo with machine-checked module walls, CI with every gate wired and proven able to fail, one-command local stack, tokens + first two components.

Deliverables. T0.1 skeleton+boundaries (DONE) - T0.2 docker/env/README - T0.3 CI gates + violating fixtures - T0.4 contracts+migrations+append-only pattern - T0.5 job table + SSE walking skeleton - T0.6 tokens + Evidence Chip + Gate Badge - T0.7 paperwork

SESSION PROMPT (after the Standing Header)

LANE: infra. Execute PLAN.md Phase 0, task .
Read first: ARCHITECTURE.md S1, S4, S13; PLAN Phase 0;
PHASE0_START_HERE.md; accepted ADR-0001.
You are building enforcement before code. Every gate you wire must be proven by a deliberately-violating fixture that FAILS CI with a named, readable error. Nothing product-shaped. No real model calls (gateway stub only). Read-back must list the gates and the fixture proving each.

Before you move on. All Section-12 validation boxes in PHASE0_START_HERE ticked with evidence links; 30-min clean-machine test passed; concierge outreach SENT (DoD item 7).

P1 - Identity, Workspace & Tenancy

Window Wk 2-5
Model Opus 4.8 (security lane) + second-family review

Objective. Auth, workspaces, and database-enforced tenant isolation (RLS) before any research data exists. Event-log substrate and security audit stream.

Deliverables. Passkey/OAuth/email auth - workspace + membership - RLS on every tenant table with negative tests - append-only event log + security audit stream - deletion/export scaffolds

SESSION PROMPT (after the Standing Header)

LANE: infra/identity. Execute Phase 1, task .
Read first: ARCHITECTURE.md S6 (RLS), S10; PLAN Phase 1.
Boring auth, database-enforced tenancy. Every tenant table ships with its RLS policy AND a negative test attempting cross-tenant reads that must fail AT THE DATABASE (app-layer checks disabled in the harness).
No custom crypto, no roles beyond owner/member, no admin bypasses.
Read-back names the tables and their policies.

Before you move on. Cross-tenant read fails at the DB with app checks disabled; RLS negative suite is a permanent CI fixture; UPDATE on event log fails by privilege.

P2 - Evidence Store & Provenance

Window Wk 5-7
Model Opus 4.8

Objective. M0: immutable, typed, provenance-stamped evidence rows with REAL/ESTIMATE typing, slice queries per agent charter, staleness clocks, originals in object storage.

Deliverables. Evidence schema + append-only privileges - provenance + typing at contract level - slice query API - staleness engine - commit-time distribution summaries - content-addressed originals

SESSION PROMPT (after the Standing Header)

```
LANE: data. Execute Phase 2, task .  
Read first: ARCHITECTURE.md S6; 07-intelligence/memory-model.md (M0);  
PLAN Phase 2. Append-only by DB privilege; corrections supersede with  
lineage, never edit. Atomic commit: kill the process mid-commit in a  
test and assert ZERO rows landed. An untyped (REAL/ESTIMATE) write  
must be a compile error. Read-back states the immutability mechanism.
```

Before you move on. Same slice query twice = identical rows; mid-commit kill leaves zero rows; app role cannot mutate a committed row (privilege test).

P3 - Canon Engines

Window Wk 7-11
Model Opus 4.8 + FOUNDER SIGN-OFF (your signature at Alpha)

Objective. The PDF becomes executable law: Gatekeeper (9 disqualifiers, gates, phase-order), Calculators (waterfall, scores, budget math), Canon-as-versioned-config, bit-exact test harness.

Deliverables. Canon config + governance path - Gatekeeper functions with pass/fail/edge fixtures - Calculators (worked example EXACT) - score assembly - budget reverse-engineer - M1 reference tables - bit-exactness harness

SESSION PROMPT (after the Standing Header)

```
LANE: engines (CODEOWNER review; no I/O - build-enforced).  
Execute Phase 3, task .  
Read first: 05-appendix/business-logic-canon.md IN FULL; ERR-001 in  
98-governance/ERRATA.md; ARCHITECTURE S4 engines row.  
Pure functions over Canon config. Fixture expected-values are derived  
BY HAND from Canon formulas - NEVER copied from doc illustrations  
(ERR-001). Decimal/integer math only; floats banned in gate and money  
paths. Any Canon ambiguity = STOP and file the question. Read-back  
lists each Canon rule ID implemented and its fixture source.
```

Before you move on. Every C1-C4 rule has a named tested function; your own hand-calculation matches engine output to the cent on 3 candidates; property suite green cross-platform.

CHECKPOINT ALPHA - audit P0-P3 before touching user files

When	After P3 (solo wk ~11). Blocks Phase 4.
Focus	Append-only enforcement - engine purity - Canon fidelity (founder signs) - RLS integrity - docs-vs-code conformance

CONNECTIVITY AUDIT PROMPT (fresh session - the auditor wrote none of this code)

ROLE: Auditor (fresh session, Fable 5 or Opus 4.8; you wrote none of this code). Run the CHECKPOINT ALPHA connectivity audit over P0-P3.

1. TRACE one synthetic candidate end to end: typed evidence rows in (P2) -> Gatekeeper screen -> Calculator waterfall (P3). Verify every number in the output traces to an input row, and that identical input yields byte-identical output (NFR-1.2).
2. ATTACK the walls: attempt a cross-tenant read with app checks off; attempt UPDATE on evidence and event log; attempt an engines import of node:fs and of the gateway. All four must fail with named errors.
3. RECONCILE: run every violating fixture branch; confirm each gate still fails. Diff module READMEs vs actual imports (drift check).
4. VERIFY Canon fidelity: pick 5 random Canon rules; confirm the engine fixture derives from the Canon formula, not an illustration.
5. REPORT: findings ranked by severity, each citing the document it violates. Zero findings is a claim you must defend, not a default.

Exit criteria. All seven PLAN S4 reviews pass; founder has personally signed engine fidelity; risk registers reviewed; remediation scheduled BEFORE new feature work.

P4 - Import Pipeline

Window	Wk 11-17 (parallel with P5)
Model	Sonnet 5 (Opus 4.8 for atomic commit + quarantine logic)

Objective. The epistemic border: upload tickets -> streaming parse -> fingerprinting -> mapping -> validation flags -> atomic commit or full quarantine. The R1 drift alarm.

Deliverables. Upload tickets + direct-to-storage - streaming CSV core + hostile corpus - fingerprint registry + drift telemetry - mapping UI + confidence - parse reports + 4 exact quarantine scenarios - Xray structured capture - embeddings at ingestion

SESSION PROMPT (after the Standing Header)

LANE: intake. Execute Phase 4, task .
Read first: ARCHITECTURE S7; 07-intelligence/domain-reasoning.md S1-S2;
06-design/screens/discover-import.md; modals spec S5.
Five-stage border: identify (INTENT-BLIND) -> parse (streaming, bounded memory) -> map (confidence + human confirm) -> validate (FLAG, never fix)
-> atomic commit or zero-row quarantine. Hostile files are the norm:
test malformed, truncated, wrong-type, formula-laden, 100k-row files.
Read-back names the four quarantine scenarios you will reproduce.

Before you move on. Design partners' real exports parse >=95%; 100k rows <=60s in bounded memory; zero partial ingestion possible; every rejected file teaches the fix.

P5 - Onboarding & App Shell

Window Wk 11-16 (parallel with P4)
Model Sonnet 5

Objective. Capital-first onboarding with the live budget reveal, the three-region app frame, Capital Meter, settings, keyboard foundation.

Deliverables. Onboarding steps 1-4 + reveal choreography - Capital Profile + change-diff - Capital Meter all states - navigator + keyboard map - settings + privacy controls

SESSION PROMPT (after the Standing Header)

LANE: UX. Execute Phase 5, task .
Read first: 06-design/screens/onboarding.md; 08-ui/screens/auth-and-onboarding.md + shell-and-pipeline.md SA; 08-ui/tokens.md.
Tokens ONLY - a raw hex/px value is a lint failure. The reveal renders Calculator output verbatim: NO frontend arithmetic (Prohibition 3 extends to the client). Reduced-motion parity for every animation.
Build to spec; park improvement ideas in PARKED-IDEAS.md.

Before you move on. A design partner completes onboarding unaided; reveal figures are Evidence Chips backed by Calculator traces; sub-\$3k copy exact per spec.

P6 - Research Pipeline & Object Model

Window Wk 17-22
Model Sonnet 5 (Opus 4.8 for FK law + status machine)

Objective. Market/Candidate objects with the structural iceberg law (FK), the board, kill/park ceremonies with receipts, kill log + resurfacing, Market View v1 with the Waterline.

Deliverables. Object model + FK law - board + read model + SSE - ceremonies (weight motion, capital-protected receipts) - kill log + resurfacing - Market View + Three Checks - Library + Canon reader - Morning Line v1 (deterministic, per ERR-003)

SESSION PROMPT (after the Standing Header)

LANE: UX/feature. Execute Phase 6, task .
Read first: 06-design/screens/pipeline.md + market-view.md; 08-ui shell-and-pipeline SB; component lexicon; ERR-003, ERR-004.
The Candidate->Market FK is the iceberg law: write the test proving scoring a marketless candidate is IMPOSSIBLE. Ceremonies frame-by-frame per spec - kill has dignity and a receipt; board drag stays disabled; kill/park reasons are structured fields. Morning Line v1 is a plain deterministic template - it must NOT simulate committee voice (ERR-003).

Before you move on. Marketless scoring impossible (FK test); kill ceremony matches spec frame-by-frame; resurfacing warns on re-approach of a killed keyword; Waterline legible in grayscale print.

P7 - Discovery Sessions

Window Wk 22-26

Model Sonnet 5

Objective. The four route runners: Canon filter presets with rationale, deviation flagging, the Screening theater over real Gatekeeper results, triage, session records. Completes the ZERO-AI sellable product.

Deliverables. Route runner + filter rationale - deviation persistence chain - Screening UX - triage + session close + records - free-signals + reverse-ASIN variants - post-mortem onboarding hook

SESSION PROMPT (after the Standing Header)

LANE: UX/feature. Execute Phase 7, task .
Read first: 06-design/screens/discover-import.md; 08-ui/screens/research-and-upload.md; PLAN Phase 7.
Filter rationale strings come from Canon config (no unexplained numbers). Deviations annotate every downstream artifact forever. The Screening paces ALREADY-COMPUTED gate results (120ms/row is UI-side; never fake engine latency). Every session ends in status changes plus exactly one next action.

Before you move on. First-time user completes keyword-first route unaided (5 usability tests); end-to-end session golden path in CI; session records show routes/filters/deviations/kills exactly.

CHECKPOINT BRAVO - audit P4-P7 = MILESTONE B Go/No-Go

When After P7 (solo wk ~26). Gates ALL AI-phase spend.

Focus The complete zero-AI product - business evidence - Layer-1 integrity under real user data

CONNECTIVITY AUDIT PROMPT (fresh session - the auditor wrote none of this code)

ROLE: Auditor (fresh session, Fable 5). CHECKPOINT BRAVO audit, P4-P7, plus Milestone B evidence review.

1. TRACE the full Layer-1 loop with a REAL Helium 10 export: upload -> parse report -> evidence rows -> Screening kills (rules cited) -> promoted candidates -> market built -> Three Checks -> board.
At every hop verify provenance survives (click chips to rows).
2. CONNECTIVITY: change the capital profile; verify downstream flags fire (budget drift). Supersede an evidence row; verify staleness propagates to every dependent surface within 1s.
3. CEREMONY QA: kill a candidate slowly. Is it dignified? Receipt correct? Kill log + resurfacing working?
4. REGRESSION: full gate battery + all violating fixtures + P1 RLS suite + P2 atomicity tests still green.
5. BUSINESS GATE (founder, not AI): >=10 design partners active, kill-acceptance >=60%, concierge revenue exists. Decision recorded: proceed to AI phases / hold and harden.

Exit criteria. Seven reviews pass; Milestone B demo done live for partners; Go/No-Go decision recorded with evidence; refactoring window (1 wk) taken BEFORE P8.

P8 - Artifact Store & Claim Checker

Window Wk 24-28 (starts parallel with P7)

Model Opus 4.8 + second-family review

Objective. The crash-test wall before any car: typed claims with mandatory citations, numeric recomputation, cross-artifact consistency, the bounce protocol, invalidation graph. No AI exists yet - by design.

Deliverables. Claim/artifact schemas + typed write path - lineage graph - mechanical checks - bounce state machine (max-2, loud failure) - invalidation graph + minimal re-run sets - version/diff surfaces

SESSION PROMPT (after the Standing Header)

LANE: data/engines. Execute Phase 8, task .

Read first: 07-intelligence/verification-and-trust.md S2-S3;
memory-model M3; ARCHITECTURE S4.

You are building the wall before any car exists. A fact claim without a citation must be UNWRITABLE at the type level. Numeric claims are recomputed by Calculators at audit; off-by-a-cent bounces. There is NO bypass path, including for tests. Read-back explains how a fabricated citation becomes unwritable.

Before you move on. Hand-authored lying artifact bounces with the claim named; 3-deep lineage staleness propagates $\leq 1s$; artifacts append-only by privilege.

P9 - AI Runtime & Model Gateway

Window Wk 28-33

Model Opus 4.8; gateway config reviewed by second family

Objective. The machinery that runs minds: persisted context packages (sealed briefcases), versioned prompt templates, agent job harness, multi-provider gateway with DA family separation, degraded mode.

Deliverables. Package assembler + persistence - template system + provenance tuple - job harness + checkpoints + bounce wiring - gateway (adapters, capability classes, cost accounting, no-retention checks) - entailment verifier slot - degraded mode - replay harness on real packages

SESSION PROMPT (after the Standing Header)

LANE: infra/committee. Execute Phase 9, task .

Read first: 07-intelligence/agent-roster.md (boundaries) +
orchestration-and-scheduling.md; ARCHITECTURE S5.

Context packages are assembled server-side from charter slices and
PERSISTED BEFORE EVERY RUN (reproducibility is the audit promise).

Prove: over-slice assembly fails; same package $\Rightarrow$ same claims, twice.
Verification failures NEVER retry silently; no silent model downgrade
on verdict-path work - it queues. Gateway credentials exist only in
the agent worker pool env.

Before you move on. Null-agent round-trip reproducible; contraband test passes; cost accounting reconciles $\pm 2\%$; provider-outage drill shows honest degraded mode.

CHECKPOINT CHARLIE - audit P8-P9. HARD STOP before agents

When	After P9 (solo wk ~33). NO agent work may start until this passes.
Focus	Package boundary physics - replay determinism - the impossibility of unverified AI output

CONNECTIVITY AUDIT PROMPT (fresh session - the auditor wrote none of this code)

ROLE: Auditor (fresh session, Fable 5). CHECKPOINT CHARLIE audit over P8-P9. This is the last gate before AI output exists; audit like it.

1. PROVE THE WALL: attempt every bypass you can construct - publish an artifact skipping the Claim Checker; cite a nonexistent row; cite a real row that does not support the claim; emit a numeric claim 1 cent off; write from a test helper. Every path must fail closed.
2. PROVE THE BRIEFCASE: build a context package exceeding a charter slice; assembly must fail. Run the null agent twice on one package; claims must be identical. Verify packages persist BEFORE runs.
3. PROVE THE GATEWAY: stub-only in CI; keys absent from all non-agent envs; degraded-mode drill end to end; cost meter vs invoice.
4. CONNECTIVITY: invalidation from P2 evidence supersession reaches P8 artifacts and computes the minimal re-run set correctly.
5. REPORT with severity + document citations. A pass here is a license to build minds; treat false confidence as the enemy.

Exit criteria. Seven reviews + the golden-session harness running real packages on the stub. Founder confirms: verification exists, therefore agents may.

P10 - The Four Analyst Agents

Window	Wk 33-44 (the long middle)
Model	Fable 5 / Opus 4.8 for templates; Sonnet 5 for contract tests + tab UX; reviews second-family

Objective. Market -> Customer -> Economics -> Risk, each: output contract, charter tests, template, tab UX, golden tolerances - behind a per-agent flag to partners.

Deliverables. Per agent x4: contract + fixtures, template + charter tests, tab UX + progressive render, golden tolerances + bench seeds, flag rollout + bounce telemetry. Plus Differentiation workbench, exception dialog, stale-dot rail.

SESSION PROMPT (after the Standing Header)

LANE: agent-template (ONE agent per session; UX in separate sessions).
Execute Phase 10, task for .
Read first: agent-roster.md charter; domain-reasoning.md
matching section; the tab spec in 08-ui candidate-and-evidence.md.
The agent judges; machines count and compute - a Customer Analyst
count or an Economics computed number is a contract-test FAILURE.
Every template change reruns goldens (gate G4). Never touch another
agent or the UX lane in this session.

Before you move on. Each agent passes goldens at tolerance; counting fidelity 100%; Economics emits zero computed numbers; partners see artifacts that survive their scrutiny.

P11 - Committee Protocol & Verdict

Window Wk 44-51

Model Fable 5 / Opus 4.8; DA on second family BY ARCHITECTURE; billing paths Opus + human line-read

Objective. Devil's Advocate, Chair, five movements, credits ledger, the memo with its ceremony (deliberation -> pause -> seal). Milestone C: first unaided verdict.

Deliverables. DA + second-family routing - Chair synthesis + adjudication - movement chain + pinning + abort - credits ledger + refunds - confidence assembly (heuristic, labeled) - memo store + document UX + ceremony - cohort ranking

SESSION PROMPT (after the Standing Header)

LANE: committee. Execute Phase 11, task .
Read first: 07-intelligence/committee-protocol.md IN FULL; verdict memo specs (06-design + 08-ui screens).
Five movements as a checkpointed chain; the record PINS at convening; mid-verdict evidence aborts cleanly. The DA runs BEFORE the Chair drafts (adversary before advocate) and on a different model family.
A memo with an unadjudicated objection is unpublishable - machine-enforced. Credit debit joins the convening transaction; failures auto-refund. The pause and the stamp are product law.

Before you move on. Golden candidate reproduces the agreed memo at tolerance; induced failure at every movement refunds correctly; a NO memo carries 'what would change this'; practitioner bench signs 5 memos.

CHECKPOINT DELTA - audit P10-P11 = MILESTONE C

When After P11 (solo wk ~51).

Focus Verdict integrity end to end - money correctness - sycophancy resistance - memo quality by humans

CONNECTIVITY AUDIT PROMPT (fresh session - the auditor wrote none of this code)

ROLE: Auditor (fresh session, Fable 5) + FOUNDER for sections 4-5.
CHECKPOINT DELTA audit over P10-P11.

1. TRACE a full verdict: real export -> evidence -> four artifacts -> Counter-Brief -> adjudications -> sealed memo. Replay 5 random memo figures to their CSV rows. Verify the version tuple stamps.
2. ATTACK: pressure the committee to pass a gated product (sycophancy bench live); attempt YES past a failed gate (path must not exist); attempt convening with a stale artifact (must refuse).
3. MONEY: induce failures at every movement; ledger must reconcile to the cent; refunds automatic; no double-debit under retries.
4. HUMAN QA (founder): watch the full ceremony as theater-critic - deliberation narration, THE PAUSE, the stamp. Read 5 memos cold: would you hand them to your own family with your money at stake?
5. Bench dashboard review: bounce rates, golden pass rates, cost per verdict vs the NFR-6.1 envelope at target pricing.

Exit criteria. Seven reviews; first unaided verdict exists; a partner acts on one (samples ordered or NO thanked); cost-per-verdict inside margin envelope. Refactoring window before P12.

P12 - Committee Conversation

Window Wk 51-55 (parallel with P13)

Model Opus 4.8 (voice + challenge logic); Sonnet 5 for panel UX

Objective. The Chair speaks: grounded panel, five Challenge Mode triggers, the argument loop that concedes only to evidence, banned utterances enforced, Morning Line v2.

Deliverables. Panel + grounding + chip-verified streaming - trigger detectors + interjections - argument loop + resolution states - scripted moments + caps + Morning Line v2

SESSION PROMPT (after the Standing Header)

LANE: committee/UX (separate sessions). Execute Phase 12, task .
Read first: 06-design/committee-conversation.md IN FULL; 08-ui candidate-and-evidence SC.
One voice (the Chair); every factual chat claim passes the same checker path as artifacts; numbers never stream half-formed. The five challenge triggers are deterministic detectors. Banned utterances ('this could work', unqualified agreement) lint at CI AND runtime.
The committee may never fold to pressure alone - concede only to evidence.

Before you move on. Sycophancy bench passes on the argument loop; you personally fail to break it in an hour; zero false-positive triggers on the golden regression day.

P13 - Reports, Handoff & Sharing

Window Wk 51-56 (parallel with P12)

Model Sonnet 5

Objective. Print-grade documents: memo PDF with chips-to-endnotes, all artifact report forms, the strategy-stripped Spec Sheet, Launch Handoff Pack, revocable share links.

Deliverables. Render pipeline + endnotes - artifact report forms - spec sheet + exclusion proof - handoff assembly - share links + revocation - account summary

SESSION PROMPT (after the Standing Header)

LANE: feature. Execute Phase 13, task .
Read first: 06-design/reports-and-artifacts.md; 08-ui/screens/verdict-and-reports.md.
The renderer reads ONLY published artifacts - by construction it cannot fabricate. The supplier Spec Sheet must contain ZERO market/keyword/margin data: write the schema-level exclusion test. Print discipline: grayscale, A4+Letter, no clipped tables. Share links private-by-default, revocable, no evidence-drawer access.

Before you move on. Memo prints beautifully on real paper in grayscale; spec-sheet strip proven by test; revocation immediate.

P14 - Hardening & Private Beta

Window Wk 56-61

Model Sonnet 5 for fixes; Opus 4.8 + EXTERNAL HUMANS for security; founder runs the fake Sev-1

Objective. Production grade: load to budgets, DR restore drills with measured RTO/RPO, external security review, a11y + state sweeps, 20+ beta users, incident playbooks rehearsed.

Deliverables. Perf/load - DR + restore drills - external security review + fixes - state-inventory + a11y sweeps - beta onboarding + metrics - incident playbooks incl. the honesty step

SESSION PROMPT (after the Standing Header)

LANE: infra/QA. Execute Phase 14, task .
Read first: 98-governance/06-QUALITY-SYSTEM.md; PLAN Phase 14;
ARCHITECTURE S12.
A backup that has not been restored is a rumor: restore with your own hands, measure RTO/RPO. Run one fake Sev-1 end to end INCLUDING the product-honesty step. Onboard beta slower than you can support.
Freeze mindset: no new features in this phase.

Before you move on. Four clean weeks: zero integrity Sev-1s; budgets green under beta load; drills documented; external review remediated.

P15 - Commercialization

Window Wk 61-64

Model Opus 4.8 (money) + human line-read of every billing path

Objective. Tiers live (Scout/Committee/Operator), 14-day trial with 1 credit, top-ups, processor integration (tokens only), candid limit states, reconciliation.

Deliverables. Plans + limits + candid states - processor + webhooks + reconciliation - trial + credit flows - invoices + tax hooks

SESSION PROMPT (after the Standing Header)

LANE: billing. Execute Phase 15, task .
Read first: 03-business/monetization-and-pricing.md; 08-ui library-settings-billing SC; ARCHITECTURE S4 billing row.
Ledger-style credits (append-only, every entry permanent). Card data NEVER touches us - processor-hosted fields only. Write the induced-failure reconciliation suite: break payments on purpose, books stay true. Prohibition 6 test: engines unreachable by plan config - gate integrity identical across tiers.

Before you move on. Reconciliation clean over the induced-failure suite; Prohibition 6 contract test green; pricing said out loud to 10 humans before the page ships.

P16 - GA Launch & Operate

Window Wk 64-67
Model Sonnet 5; humans hold ship authority

Objective. The door opens: production floors, status page, launch runbook executed as rehearsal first, ECHO passed, GA.

Deliverables. Production scaling floors - status page + on-call - launch runbook rehearsed - ECHO + Go/No-Go - launch - post-launch triage cadence

SESSION PROMPT (after the Standing Header)

LANE: infra. Execute Phase 16, task .
Read first: ARCHITECTURE S13; 98-governance/08-RELEASE-GOVERNANCE.md.
Nothing new ships launch week. The runbook is executed as a full rehearsal before the real day. Rollback paths named in the release notes. AI agents never hold ship authority - a human ships.

Before you move on. Cold production walkthrough: trial -> pay -> verdict end to end; ECHO signed; a stranger pays.

CHECKPOINT ECHO - final audit P12-P16 before GA

When Before GA (solo wk ~66). The last gate.
Focus Whole-system integrity - launch readiness - the candor surfaces

CONNECTIVITY AUDIT PROMPT (fresh session - the auditor wrote none of this code)

ROLE: Auditor (fresh session, Fable 5) + founder for section 4.
CHECKPOINT ECHO - the pre-GA audit over the ENTIRE system.

1. FULL-JOURNEY TRACE on production config: signup -> onboarding -> import -> screening -> market -> phases -> verdict -> memo PDF -> share link -> billing. Every chip clicks through to its row.
2. REGRESSION TOTALITY: entire gate battery, all violating fixtures, RLS suite, atomicity suite, golden sessions, all benches, print snapshots, reconciliation suite. 100% green or no launch.
3. OPERATIONS: status page live; on-call rota staffed; DR evidence fresh (<90 days); rollback rehearsed on the launch build itself.
4. CANDOR CHECK (founder): read the landing page, trial copy, memo disclaimer, and pricing page in one sitting. Is every claim true TODAY? Pre-calibration confidence language prominent? The honest NO visible in marketing?
5. Go/No-Go per PLAN 5.5 recorded with evidence links.

Exit criteria. Milestone E Go/No-Go signed. Then: launch, watch, write everything down, rest one week - and read Part 12 of the Founder Execution Guide.

5 - Progress Tracker

Tick with a pen. A phase is done when its exit checks hold with evidence links, not when the code compiles.

Item	Done	Date	Evidence link
P0 - Repository & Foundation	☐		
PM-1: 10 paid concierge verdicts (parallel, wk 1-9)	☐		
P1 - Identity, Workspace & Tenancy	☐		
P2 - Evidence Store & Provenance	☐		
P3 - Canon Engines	☐		
CHECKPOINT ALPHA - audit P0-P3 before touching user files	☐		
P4 - Import Pipeline	☐		
P5 - Onboarding & App Shell	☐		
P6 - Research Pipeline & Object Model	☐		
P7 - Discovery Sessions	☐		
CHECKPOINT BRAVO - audit P4-P7 = MILESTONE B Go/No-Go	☐		
P8 - Artifact Store & Claim Checker	☐		
P9 - AI Runtime & Model Gateway	☐		
CHECKPOINT CHARLIE - audit P8-P9. HARD STOP before agents	☐		
P10 - The Four Analyst Agents	☐		
P11 - Committee Protocol & Verdict	☐		
CHECKPOINT DELTA - audit P10-P11 = MILESTONE C	☐		
P12 - Committee Conversation	☐		
P13 - Reports, Handoff & Sharing	☐		
P14 - Hardening & Private Beta	☐		
P15 - Commercialization	☐		
P16 - GA Launch & Operate	☐		
CHECKPOINT ECHO - final audit P12-P16 before GA	☐		

Close this playbook. Open PLAN.md. Execute the next task.